

Amal Mehta

SF Bay Area · amal.mehta@gmail.com · amalmehta.github.io/PortfolioWebsite · linkedin.com/in/amal-mehta-264413a9

Research engineer with four years of applied deep learning across the stack: generative models and vision (GANs, CNN segmentation, detection) and deep RL and meta-RL, from data pipelines and multi-GPU training to policies deployed on physical robots. Prior neuroimaging research at UCSF; long-standing interest in how the brain learns.

Experience

Independent ML & Robotics Research

Sept 2025 – Present

Self-directed projects; details at amalmehta.github.io/PortfolioWebsite

San Francisco Bay Area

- **Neural decoding:** 2D hand velocity from motor/premotor cortex spikes (NLB MC_Maze); GRU decoder outperformed a Kalman filter on held-out trials (R^2 0.75 vs. 0.55).
- **Post-training:** SFT then GRPO on a small transformer policy in a key-and-door gridworld maze; GRPO lifted held-out success from 71% to 85%, where RL from scratch failed.
- **Neuroimaging:** macOS tool for 3D brain MRI; gray/white matter and CSF segmentation, plus EEG electrode placement.
- **Pose estimation:** real-time sport technique analysis on MediaPipe 3D pose from a webcam; scores reps against one reference rep.
- Member, Sensorimotor AI journal club, Redwood Center for Theoretical Neuroscience, UC Berkeley.

Deep Learning Scientist

Sept 2021 – Aug 2025

PhysicsAI (DARPA-funded)

Pacifica, CA (remote)

- Formulated synthetic-data generation as RL: policies (TD3, PPO, DDPG) set image-generator parameters, rewarded by detector performance on real drone imagery, closing the sim-to-real gap without hand-tuned rendering.
- Built a PEARL meta-RL agent (probabilistic task embeddings over TD3) for fast few-shot adaptation across generator tasks.
- Implemented domain adaptation from scratch to hold detector performance in the low-real-label, high-synthetic-data regime.
- Built GAN-based synthesis pipelines (DCGAN, InfoGAN, FineGAN) producing training data for downstream detection models.

Deep RL Researcher

Aug 2023 – Mar 2024

Hybrid Robotics Lab, UC Berkeley — adv. Prof. Koushil Sreenath

Berkeley, CA

- Trained a single multi-stage PPO policy in IsaacGym for a Unitree Go1 quadruped to track, approach, and shoot a soccer ball at a target, replacing the multi-policy hierarchical framework used in prior work; ~60% kicking accuracy on hardware.
- Designed the three-stage reward and transition machinery and CPG-based gait tracking; prototyped RLPD (SAC with prior data) to fine-tune the sim-trained policy on real robot data.

Research Assistant

June 2020 – Dec 2022

Video & Image Processing Lab, BAIR — adv. Prof. Avideh Zakhori

UC Berkeley

- Developed a deep CNN segmentation pipeline for invasive melanoma in large-scale histopathological images; published at SPIE Medical Imaging 2023; Stowell-Orbison Award Poster Session at USCAP 2022.

Researcher

June 2020 – Aug 2020

Abbasi-Asl Lab, UCSF — adv. Prof. Reza Abbasi-Asl

San Francisco, CA

- Built empirical mode decomposition (EMD) pipelines for 2D and 3D brain image registration.

Education

University of California, Berkeley

Aug 2017 – Aug 2021

B.A. Computer Science — Regents Scholar

Berkeley, CA

Stanford University

Aug 2021 – May 2023

Graduate coursework (non-degree) — Autonomous Systems & Artificial Intelligence

- CS231N (Deep Learning for Vision), CS330 (Meta-Learning), AA271A/B (Autonomy I & II), AA273 (State Estimation).
- CS330 project: few-shot imitation learning; implemented RL² and inverse soft Q-learning in Meta-World.

Technical Skills

Languages	Python, C++, C
ML / RL	PyTorch, TensorFlow, stable-baselines3, Gym, Meta-World, IsaacGym, OpenCV
Robotics	ROS, Unitree Go1, CPG gait control, Kalman filtering
Tools	Git, GitHub Actions CI, Linux, Docker, CUDA/multi-GPU, W&B, TensorBoard

Publications

1. Shah, A., **Mehta, A.**, Wang, M., Neumann, N., McCalmont, T., Zakhori, A. “Deep Learning Segmentation for Invasive Melanoma.” *SPIE Medical Imaging*, 2023.
2. Neumann, N., Wang, M., **Mehta, A.**, et al. “Quantifying Invasive Melanoma Volume by Deep Learning Segmentation at the Pixel-Level.” *USCAP Annual Meeting*, 2022.